

DFT & FFT Implementation with CUDA

Group 17

A59023194 Ya-Chi Liao

A59026653 Tsung-Hsiang Ma

Introduction

- **Importance of DFT and FFT in Modern Applications**
 - Widely used in signal processing, image processing, communications, and scientific simulations.
 - Essential for efficient frequency domain analysis.
- **Performance Bottleneck in Traditional Approaches**
 - Computationally intensive for large datasets.
 - Serial implementations struggle with real-time processing demands.
- **Growing Need for High-Performance Solutions**
 - Explosion of big data in AI, IoT, and multimedia applications.
 - Real-time requirements in fields like radar, medical imaging, and financial modeling.
- **Role of GPUs and CUDA**
 - GPUs offer massive parallelism ideal for computationally heavy tasks.
 - CUDA programming allows fine-grained control of GPU resources.

Objective

- **Optimize Computational Performance:**

Exploit parallel processing capabilities of GPUs to reduce computation time.

- **Resource Efficiency:**

Minimize memory usage and optimize GPU resource utilization.

Methodologies

- **Discrete Fourier Transform (DFT)**
 - Benchmark
 - Acceleration with CUDA global memory
 - Acceleration with CUDA shared memory
- **Fast Fourier Transform (FFT)**
 - Benchmark
 - Acceleration with CUDA global memory
 - Implementation with CuFFT for reference

DFT Introduction

- **Discrete Fourier Transform (DFT):**

- The DFT converts a finite sequence of equally spaced samples of a function into a sequence of coefficients of a finite combination of complex sinusoids.
- Direct computation of DFT has $O(n^2)$ complexity, which can be significantly reduced using the Fast Fourier Transform algorithm with $O(N \log N)$ complexity.
- $X[k] = \sum x[n] \cdot e^{(-2\pi knj/N)}$, $k, n=0, 1, \dots, N-1$

DFT Implementation

- Thread Organization
 - Assign one CUDA thread for computing each output frequency component $X[k]$.
 - Use shared memory to minimize global memory access during computation.
- Kernel Design
 - Each thread computes $X[k]$ by iterating over all $x[n]$ values, performing the sum of the products $x[n] \cdot e^{(-2\pi knj/N)}$.
- Optimization
 - Use Coalesced Memory Access to fetch $x[n]$ from global memory.
 - Use Fast Math Operations to compute sine and cosine efficiently.

DFT Benchmark

- This benchmark serves as a reference for comparing speed and accuracy with CUDA-accelerated implementation.
- $X[k] = \sum x[n] \cdot e^{(-2\pi knj/N)}$, $k, n=0, 1, \dots, N-1$

```
void computedDFT_CPU(const float* input_real, const float* input_imag,
float* output_real, float* output_imag, int N) {
    for (int k = 0; k < N; ++k) { // Iterate over each frequency bin
        double sum_real = 0.0;
        double sum_imag = 0.0;

        for (int n = 0; n < N; ++n) { // Sum over the input signal
            double angle = -2.0 * M_PI * k * n / N;
            double cos_val = cos(angle);
            double sin_val = sin(angle);

            sum_real += input_real[n] * cos_val - input_imag[n] * sin_val;
            sum_imag += input_real[n] * sin_val + input_imag[n] * cos_val;
        }

        output_real[k] = sum_real;
        output_imag[k] = sum_imag;
    }
}
```

Global Memory DFT

- The input signal consists of N complex numbers, where each number has a real and imaginary component.
- The output consist of
 - $\text{real}(X[k]) = \sum x[n] \cdot \cos(2\pi kn/N)$
 - $\text{img}(X[k]) = \sum x[n] \cdot \sin(2\pi kn/N)$
- Compute $X[k]$ for $k = 1, 2, \dots, N-1$

Shared Memory DFT

- Input Tiling:
 - Divide the input data into tiles that fit into the shared memory of each block.
- Shared Memory Allocation:
 - Allocate shared memory dynamically within the kernel to store a portion of the input data for each thread block.
- Loading Data into Shared Memory
- Synchronization
- Local Computation
- Iterate Over Tiles

Implementing DFT with shared memory is good but inefficient for large N . FFT algorithms are more practical in real-world applications.

Result

N	16	256	1024
Benchmark	37.79 us	936 us	3.70 ms
Global Mem	37.82 us	934 us	3.71 ms
Shared Mem	38.50 us	917 us	3.52 ms

FFT Introduction

- **Fast Fourier Transform (FFT):**

- An efficient algorithm to compute the Discrete Fourier Transform (DFT).
- Converts a signal from the time domain to the frequency domain.
- Reduces computational complexity from $O(n^2)$ (naive DFT) to $O(n \log n)$

- **FFT Implementation:**

- Bit-Reversal Reordering
 - Reorganizes input data for efficient memory access.
- Iterative Cooley-Tukey Algorithm
 - Processes data in iterative stages.
 - Combines results using "butterfly" operations for each pair of input points.
 - Utilizes twiddle factors to compute complex multiplications.

FFT Benchmark

- This benchmark serves as a reference for comparing speed and accuracy with CUDA-accelerated implementation.
- **Code Structure:**
 - Bit-reversal reordering
 - Iterative Cooley-Tukey FFT algorithm:
 - Radix-2 Iterative Implementation
 - Twiddle Factor Precomputation

Global Memory FFT

- **Bit-Reversal Reordering**
 - Parallel Execution: Each thread processes one index, calculating its bit-reversed counterpart independently.
 - GPU Threads: Kernel `bit_reversal_kernel` launches N threads for N -point FFT.
 - Reduced Time Complexity: Instead of serial reordering, all elements are processed concurrently.
- **Twiddle Factor Computation**
 - Dynamic Calculation in Each Thread: Each thread computes its unique twiddle factor during the "butterfly" operation.
 - Thread Independence: Each thread works on its portion of the data, eliminating dependencies and enabling parallelism.
- **Iterative Cooley-Tukey**
 - Parallel Butterfly Operations: Each thread handles one butterfly computation (even-odd pair).
 - Efficient Memory Access: Reordered input (bit-reversal) ensures coalesced memory access patterns.

Result

N	16	256	1024
Benchmark	41.7 us	7.74 ms	37.1 ms
Global Mem	36.28 us	76.45 us	112.37 us
CuFFT(for reference)	10.21 us	10.21 us	10.27 us

Conclusion

- **Performance Trends:**
 - Shared memory optimizations slightly reduce computation time for DFT, especially for larger N . However, gains are marginal.
- **FFT Advantage:**
 - FFT with GPU is significantly faster than DFT and FFT with CPU due to its optimized $O(N\log N)$ complexity and the capability of being parallel computing.
- **Applicability:**
 - DFT with CUDA is educational, but FFT is the practical choice for efficient frequency-domain transformations in GPU applications.

Reference

- Govindaraju, N. K., Lloyd, B., Dotsenko, Y., Smith, B., & Manferdelli, J. (2008, November). High performance discrete Fourier transforms on graphics processors. In SC'08: Proceedings of the 2008 ACM/IEEE conference on Supercomputing (pp. 1-12). IEEE.
- Nvidia, CuFFT Library, https://docs.nvidia.com/cuda/pdf/CUFFT_Library.pdf